

CSE5306 PA3 — Fault-Tolerant Extensions Using 2PC and Raft

Faisal Ahmad
Student ID: 1002239354

Project Repository

[Link](#)

1 Introduction

This report documents the design and implementation of Two-Phase Commit (2PC) and Raft consensus algorithms added to the base distributed polling system from CSE5306 Project 2 (Made by Group-3). The work extends a REST-based polling backend with two independent fault-tolerance mechanisms.

[Link to Original Base Project Repository](#)

2 Original Base Project

Summary of base project:

- REST API implemented using Hono (Node.js)
- Stateless backend containers (api1, api2)
- PostgreSQL single database instance
- Nginx load balancer in front of APIs
- Standard CRUD for polls and votes

The base system had **no fault-tolerant write mechanism**. Both 2PC and Raft implementations extend the **CreatePoll** path only.

3 Two-Phase Commit (Q2)

3.1 System Architecture

The 2PC version adds:

- One Coordinator (inside `api/`)
- Four Participant Nodes (`two_pc_participant/`)
- gRPC communication via `two_pc.proto`

Data flow for `/polls`:

1. API receives poll creation request
2. API starts a 2PC transaction:
 - Sends `RequestVote` to all participants
 - If all vote YES: send `SendDecision(COMMIT)`
 - Else: `SendDecision(ABORT)`
3. API writes poll to DB only after COMMIT

This ensures atomicity.

4 Raft Leader Election (Q3)

4.1 Architecture

Implemented in:

`base_rest_raft/raft_node/src/server.ts`

Each of 5 Raft nodes:

- Starts as follower
- Waits randomized 1.5–3.0s election timeout
- If no heartbeat: becomes candidate
- Increments term
- Votes for itself
- Broadcasts `RequestVote` to peers
- If receives majority: becomes leader
- Sends heartbeat every 1s

The API communicates with Raft using the entry node.

5 Raft Log Replication (Q4)

5.1 Operation Flow

When a client creates a poll:

1. API sends `HandleClient` RPC to Raft
2. If caller is not leader: leader redirect
3. Leader appends log entry (`CREATE_POLL`)
4. Leader sends `AppendEntries` (log + commitIndex)
5. Followers overwrite log with leader's log and apply committed entries
6. Leader updates commit index immediately (assignment simplification)
7. API writes poll to DB after commit

This satisfies Q4 requirements.

6 Test Cases (Q5)

All logs are stored in:

`base_rest_raft/logs/`

Test Case 1 — Normal Operations

- Start cluster normally
- Submit poll
- Leader receives operation, replicates, commits

```
[TC1] Normal Raft operation: leader election + create poll
[TC1] Log file: logs\test1_normal_20251123_173639.log

[TC1] Creating poll via load balancer on :3005...
curl -X POST http://localhost:3005/polls -H "Content-Type: application/json" -d "{\"question\":\"TC1: normal?\", \"options\":[\"yes\", \"no\"]}"
  % Total    % Received % Xferd  Average Speed   Time     Time   Current
               Dload  Upload Total Spent    Left  Speed
0          0      0      0      0      0      0 --::-- --::-- --::-- 0
100  180  100  130  100  50  2436  937 --::-- --::-- --::-- 3396
{"id":26,"question":"TC1: normal?","options":["yes","no"],"isActive":true,"createdAt":"2025-11-23T23:36:39.616Z","raftLogIndex":2}
[TC1] Done. Check logs for leader election and log application.
[TC1] Now you can run:
  docker compose logs api1 --tail=40
  docker compose logs raft3 --tail=80
```

Test Case 2 — Leader Crash

- Identify leader using logs
- `docker compose stop raftX`
- Observe election timeout
- New leader chosen
- Submit poll → success

```
[TC2] Leader crash and re-election test
[TC2] Using assumed leader service: raft3
[TC2] Log file: logs\test2_leader_crash_20251123_173655.log

[TC2] Stopping leader container: raft3 ...
Container raft_node3 Stopping
Container raft_node3 Stopped

[TC2] Immediately trying to create a poll (may fail or be rejected)...
curl -X POST http://localhost:3005/polls ...
% Total    % Received % Xferd  Average Speed   Time   Time   Time  Current
          Dload  Upload Total Spent   Left Speed
0       0     0     0     0     0      0 --::--- --::--- --::---      0
100  204  100  142  100   62  5061  2210 --::--- --::--- --::---  7285
{"id":27,"question":"TC2: after leader crash?","options":["yes","no"],"isActive":true,"createdAt":"2025-11-23T23:36:56.278Z","raftLogIndex":3}
[TC2] Waiting 5 seconds for re-election...

[TC2] Trying again after re-election...
% Total    % Received % Xferd  Average Speed   Time   Time   Time  Current
          Dload  Upload Total Spent   Left Speed
0       0     0     0     0     0      0 --::--- --::--- --::---      0
100  202  100  141  100   61  6026  2687 --::--- --::--- --::---  8782
{"id":28,"question":"TC2: after re-election?","options":["yes","no"],"isActive":true,"createdAt":"2025-11-23T23:37:01.311Z","raftLogIndex":4}
[TC2] Now check raft node logs for a new leader being elected.
```

Test Case 3 — Follower Recovery

- Stop follower
- Leader continues operating
- Restart follower
- Leader sends AppendEntries
- Log catches up

```
[TC3] Follower crash and recovery (log catch-up)
[TC3] using follower service: raft4
[TC3] Log file: logs\test3_follower_recovery_20251123_173814.log

[TC3] Stopping follower: raft4 ...
Container raft_node4 Stopping
Container raft_node4 Stopped

[TC3] Creating 2 polls while follower is DOWN...
% Total    % Received % Xferd  Average Speed   Time   Time   Time  Current
          Dload  Upload Total Spent   Left Speed
0       0     0     0     0     0      0 --::--- --::--- --::---      0
100  202  100  141  100   61  3615  1564 --::--- --::--- --::---  5315
{"id":29,"question":"TC3: while raft4 down 1","options":["yes","no"],"isActive":true,"createdAt":"2025-11-23T23:38:15.977Z","raftLogIndex":5} % Total    % Received % Xferd  Average Speed
Time   Time   Time  Current
          Dload  Upload Total Spent   Left Speed
0       0     0     0     0     0      0 --::--- --::--- --::---      0
100  202  100  141  100   61  7856  3399 --::--- --::--- --::--- 11882
{"id":30,"question":"TC3: while raft4 down 2","options":["yes","no"],"isActive":true,"createdAt":"2025-11-23T23:38:16.032Z","raftLogIndex":6}
[TC3] Starting follower again: raft4 ...
Container raft_node4 Starting
Container raft_node4 Started

[TC3] Waiting 5 seconds for it to receive AppendEntries...

[TC3] Now manually run: docker compose logs raft4 --tail=80
[TC3] Look for: Applying log index... op=CREATE_POLL
```

Test Case 4 — Concurrent Poll Creation

- Rapid curl requests
- Log shows proper ordering and replication

```
[TC4] Multiple back-to-back client polls (log ordering test)
[TC4] Log file: logs\test4_concurrent_20251123_173835.log

[TC4] Creating poll 1 ...
% Total    % Received % Xferd  Average Speed   Time   Time   Time  Current
          Dload  Upload   Total Spent   Left Speed
0       0     0     0     0     0       0 --::--- --::--- --::--- 0
100  180  100  130  100   50  6048  2326 --::--- --::--- --::--- 8571
{"id":31,"question":"TC4: poll 1?","options":["yes","no"],"isActive":true,"createdAt":"2025-11-23T23:38:35.807Z","raftLogIndex":7}
[TC4] Creating poll 2 ...
% Total    % Received % Xferd  Average Speed   Time   Time   Time  Current
          Dload  Upload   Total Spent   Left Speed
0       0     0     0     0     0       0 --::--- --::--- --::--- 0
100  180  100  130  100   50  7429  2857 --::--- --::--- --::--- 10588
>{"id":32,"question":"TC4: poll 2?","options":["yes","no"],"isActive":true,"createdAt":"2025-11-23T23:38:35.892Z","raftLogIndex":8}
[TC4] Creating poll 3 ...
% Total    % Received % Xferd  Average Speed   Time   Time   Time  Current
          Dload  Upload   Total Spent   Left Speed
0       0     0     0     0     0       0 --::--- --::--- --::--- 0
100  180  100  130  100   50  6293  2420 --::--- --::--- --::--- 9000
>{"id":33,"question":"TC4: poll 3?","options":["yes","no"],"isActive":true,"createdAt":"2025-11-23T23:38:35.988Z","raftLogIndex":9}
[TC4] Creating poll 4 ...
% Total    % Received % Xferd  Average Speed   Time   Time   Time  Current
          Dload  Upload   Total Spent   Left Speed
0       0     0     0     0     0       0 --::--- --::--- --::--- 0
100  181  100  131  100   50  7915  3021 --::--- --::--- --::--- 11312
>{"id":34,"question":"TC4: poll 4?","options":["yes","no"],"isActive":true,"createdAt":"2025-11-23T23:38:36.079Z","raftLogIndex":10}
[TC4] Creating poll 5 ...
% Total    % Received % Xferd  Average Speed   Time   Time   Time  Current
          Dload  Upload   Total Spent   Left Speed
0       0     0     0     0     0       0 --::--- --::--- --::--- 0
100  181  100  131  100   50  8434  3219 --::--- --::--- --::--- 12066
>{"id":35,"question":"TC4: poll 5?","options":["yes","no"],"isActive":true,"createdAt":"2025-11-23T23:38:36.162Z","raftLogIndex":11}
[TC4] Now check logs on leader and follower for sequential log indexes.
[TC4] Example:
  docker compose logs raft3 --tail=80
  docker compose logs raft4 --tail=80
```

Test Case 5 — New Node Joining

- docker compose up -d raft5 after removal
- Node boots as follower
- Receives full log on heartbeat
- Catches up

```

[TCS] New node entering the system (late-joining Raft node)
[TCS] Using node: raft5
[TCS] Log file: logs\test5_new_node_join_20251123_173847.log

[TCS] Stopping raft5 so it is OUT of the cluster...
Container raft_node5 Stopping
Container raft_node5 Stopped

[TCS] Creating 3 polls while raft5 is down...
% Total % Received % Xferd Average Speed Time Time Time Current
                                         Dload Upload Total Spent Left Speed

0 0 0 0 0 0 0 -:- -:- -:- -:- -:- 0
100 62 0 0 100 62 0 303 -:- -:- -:- -:- -:- 302
100 62 0 0 100 62 0 51 0:00:01 0:00:01 -:- -:- 51
100 62 0 0 100 62 0 27 0:00:02 0:00:02 -:- -:- 27
100 62 0 0 100 62 0 19 0:00:03 0:00:03 -:- -:- 19
100 62 0 0 100 62 0 14 0:00:04 0:00:04 -:- -:- 14
100 188 100 126 100 62 25 12 0:00:05 0:00:05 -:- -:- 26
{"error":"Raft cluster did not accept operation","details":"14 UNAVAILABLE: Name resolution failed for target dns:raft5:6000"} % Total % Received % Xferd Average Speed Time Time
Time Current                                         Dload Upload Total Spent Left Speed

0 0 0 0 0 0 0 -:- -:- -:- -:- -:- 0
0 0 0 0 0 0 0 -:- -:- -:- -:- -:- 0
100 205 100 143 100 62 5142 2229 -:- -:- -:- -:- -:- 7321
{"id":36,"question":"TCS: before raft5 join 2","options":["yes","no"],"isActive":true,"createdAt":"2025-11-23T23:38:54.191Z","raftLogIndex":12} % Total % Received % Xferd Average Speed
Time Time Time Current                                         Dload Upload Total Spent Left Speed

0 0 0 0 0 0 0 -:- -:- -:- -:- -:- 0
100 205 100 143 100 62 5649 2449 -:- -:- -:- -:- -:- 8200
{"id":37,"question":"TCS: before raft5 join 3","options":["yes","no"],"isActive":true,"createdAt":"2025-11-23T23:38:54.259Z","raftLogIndex":13}
[TCS] Starting raft5 back (late join)...
Container raft_node5 Starting
Container raft_node5 Started

[TCS] Waiting 5 seconds for it to sync logs via AppendEntries...

[TCS] Now manually run: docker compose logs raft5 --tail=100
[TCS] Look for: Applying log index<1..3 op=CREATE_POLL

```

7 Directory Structure Summary

Full structure included exactly as submitted (see appendix in repo).

8 Implementation Notes and Challenges

- Election storms occurred early due to short initial timeouts
 - Fixed via correct timeout randomization
 - Database commits triggered only on successful Raft commit
 - 2PC path isolated and independent from Raft path

9 Conclusion

This project successfully implemented:

- Fully working 2PC protocol
 - Fully working Raft cluster with leader election
 - Log replication system
 - Integration with consensus mechanisms into REST stack
 - Five documented test cases verifying correctness

All Q2–Q5 requirements were satisfied.

Appendix: Log Files

Located in:

`base_rest_raft/logs/`